

# TORAX - A Fast and Differentiable Tokamak Transport Simulator in JAX

**Jonathan Citrin, Sebastian Bodenstein,  
Anushan Fernando, Ian Goodfellow, Philippe Hamel,  
Tamara Norman, Akhil Raju, Craig Donner, Federico Felici,  
Andrea Huber, David Pfau, Brendan Tracey**

Google DeepMind

**Devon Battaglia, Anna Teplukhina, Josiah Wai**

Commonwealth Fusion Systems



Photo by Khyati Trehan for Google DeepMind on Unsplash

# Fusion at Google DeepMind

Demonstrated deep reinforcement learning for pulse trajectory + magnetic controller design

Next: more physics for multi-objective optimization and controller design tasks


TORAX in this context: fast-and-accurate simulation environment for internal plasma dynamics

Article | [Open access](#) | Published: 16 February 2022


## Magnetic control of tokamak plasmas through deep reinforcement learning

[Jonas Degraeve](#), [Federico Felici](#) , [Jonas Buchli](#) , [Michael Neunert](#), [Brendan Tracey](#) , [Francesco Carpanese](#), [Timo Ewalds](#), [Roland Hafner](#), [Abbas Abdolmaleki](#), [Diego de las Casas](#), [Craig Donner](#), [Leslie Fritz](#), [Cristian Galperti](#), [Andrea Huber](#), [James Keeling](#), [Maria Tsimpoukelli](#), [Jackie Kay](#), [Antoine Merle](#), [Jean-Marc Moret](#), [Seb Noury](#), [Federico Pesamosca](#), [David Pfau](#), [Olivier Sauter](#), [Cristian Sommariva](#), ... [Martin Riedmiller](#)  [+ Show authors](#)

[Nature](#) **602**, 414–419 (2022) | [Cite this article](#)



TORAX is our new **differentiable tokamak core transport simulator** built in Python using JAX, solving for core temperatures, density, and current diffusion



# Outline

Motivations	01
TORAX overview	02
Benchmarks	03
Outlook	04

# Motivations



1


# Next generation of tokamak experiment aiming for net fusion gain

ITER: Cadarache, France, standard 5T field



Acknowledgement to ITER Organization

SPARC, Commonwealth Fusion Systems (CFS), HFS. ~12T





Acknowledgement to Commonwealth Fusion Systems



# Key issue: present-day physics models too slow for routine simulations used for experiment prediction and interpretation

Leap from present-day experiments to reactors requires leap in simulation capabilities




Reduce costs and risks with simulations for:

- Experimental preparation including inter-shot
- Performance (constrained) optimization
- Model based controller design
- Reactor design

In next generation devices, fast and accurate simulation a prerequisite for pulse design

# Integrated modelling inherently multiscale and multiphysics

Multiple orders of magnitude in spatiotemporal scales between relevant physics processes



Magnetic equilibrium

Plasma collisions

Magnetohydrodynamics

Heating and fuelling

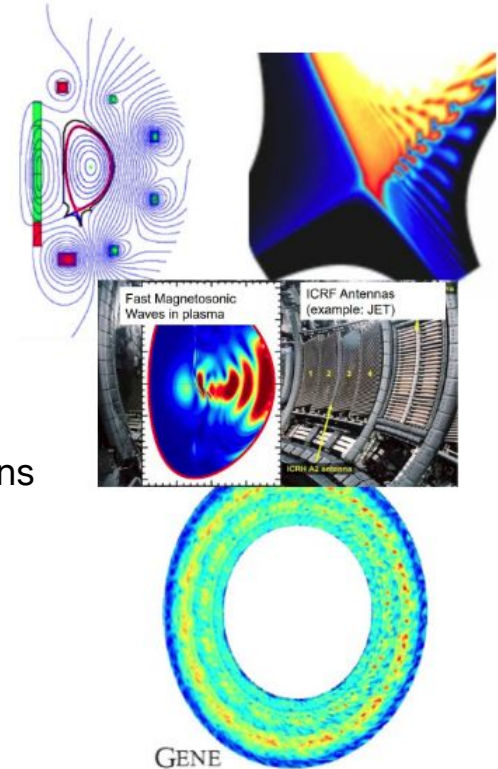
Plasma turbulence

Atomic and molecular interactions

Plasma material interaction

...


...





# Modelling fidelity and tractability hierarchy.

## Pragmatic modelling demands model reduction




Dominiski Phys. Plasmas 2018

Simulation Class	Method	Fidelity level	Compute burden	Advantages	Challenges
Direct numerical simulation	Coupling high fidelity models	High	Massively parallel and expensive (exascale)	Ultimate ground truth (numerically)	Not pragmatic for most use cases
Standard integrated modelling	Plasma transport PDEs + coupled reduced models	Variable (depends on quality of reduced models)	Hours to weeks on single compute node	Suitable for experimental interpretation and extrapolation Community workhorse. Lots of experience and models available	Legacy burden.  Tractability and accuracy are conflicting constraints
Fast integrated modelling	Plasma transport PDEs + coupled surrogate models	Variable (depends on quality of surrogates)	Faster than realtime to minutes	Suitable for optimization and controller design applications.  Surrogates can learn higher fidelity models than “standard”	Need to develop collection of learned surrogates + appropriate framework

# Integrated modelling simulation environment: Separate plasma regions in core (~1D) and edge (~2D). Different computational challenges that must be integrated

Plasma core:

1D coupled PDEs for particles, energy, rotation



$$\begin{aligned} \text{Particle density: } \frac{\partial n_s}{\partial t} + \frac{1}{r} \frac{\partial}{\partial r} (r \Gamma_s) &= S_s \\ \text{Energy: } \frac{3}{2} \frac{\partial P_s}{\partial t} + \frac{1}{r} \frac{\partial}{\partial r} (r q_s) &= Q_s \end{aligned}$$


Particle flux      Particle sources/sinks  
Heat flux      Heat sources/sinks

- Confined hot plasma region where fusion happens
- Computational bottleneck: turbulent fluxes, with highly nonlinear transport coefficients

Multiphysics: magnetic equilibrium, plasma turbulent transport, MHD, heating and fuelling

Plasma edge:

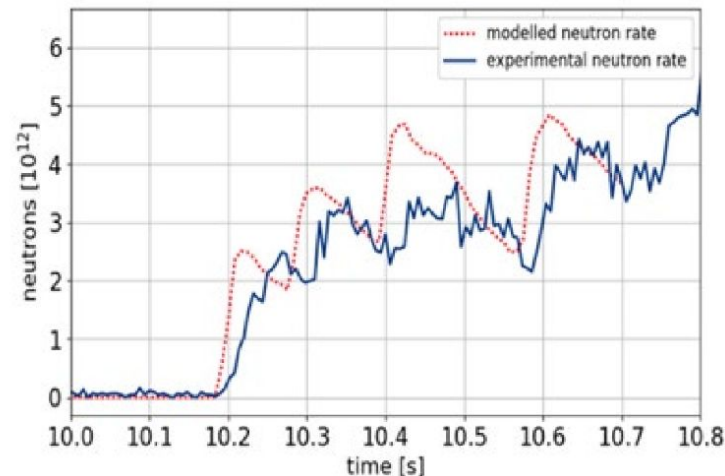

2D coupled PDEs for plasmas, 3D kinetic for neutrals



- Plasma fuelling and pumping and impact on confinement
- Impact of plasma exhaust on wall materials

Multiphysics: magnetic equilibrium, plasma + neutral transport, atomic and molecular interactions, plasma material interaction

# Example application: understanding physical mechanism of fast isotope mixing in multiple-isotope JET experiments




Marin et al Nucl. Fusion 2021


- Multiphysics simulation:
  - Physics-based turbulence model
  - Fueling with injection of frozen deuterium pellets which ablate into initial pure hydrogen plasma
  - Deuterium fusion leads to observed neutrons, measuring timescales of deuterium transport
- Acceptable sim2real gap (not trivial!) but ~1 week per simulation

# Key method for integrated modelling speed is incorporation of learned ML surrogates of physics components


Neutral Beam Injection  
and Current Drive



Magnetic  
equilibrium




Edge Transport Barriers  
(H-mode pedestal)



Meneghini NF 2017


Plasma turbulence



Lao PPCF 2022

# Intermezzo on ML-surrogates. Nonlinear regression of model databases

## Fully connected feedforward neural network (simple topology)



$x$ : Inputs, e.g. temperature and density gradients, magnetic geometry info

$y$ : Output: e.g. ion heat flux

$w$ : free weights for optimization

$$y = \sum w_j^2 g_j \left( \sum w_{i,j}^1 x_i \right)$$

With, e.g.  $g(x) = \tanh(x)$

Normally, optimize weights by minimizing cost function:  $C = \sum_N (t_N - y_N)^2 + \lambda \sum (w_{ij})^2$

$t_N$  are target values from generated dataset

$\lambda$  is the regularization factor. Critical for avoiding overfitting

- Provides an analytical formulas with analytical derivatives
- Complex nested nonlinear functions. Different versions/models can have different topology
- Autodiff is really key for incorporating ML-surrogates into differentiable simulators





# TORAX overview



2



# TORAX motivated by requirements for pulse simulation, planning, and controller design tasks

-  Fast and accurate forward modelling
-  Differentiable for accurate nonlinear PDE solvers, gradient-driven optimization and parameter identification
-  Easy incorporation of physics model ML-surrogates; higher fidelity simulation without sacrificing speed
-  Modularity for coupling within various workflows and to additional physics models.

Building on ideas developed over many years in the fusion community for tokamak scenario modeling, pulse planning, and control

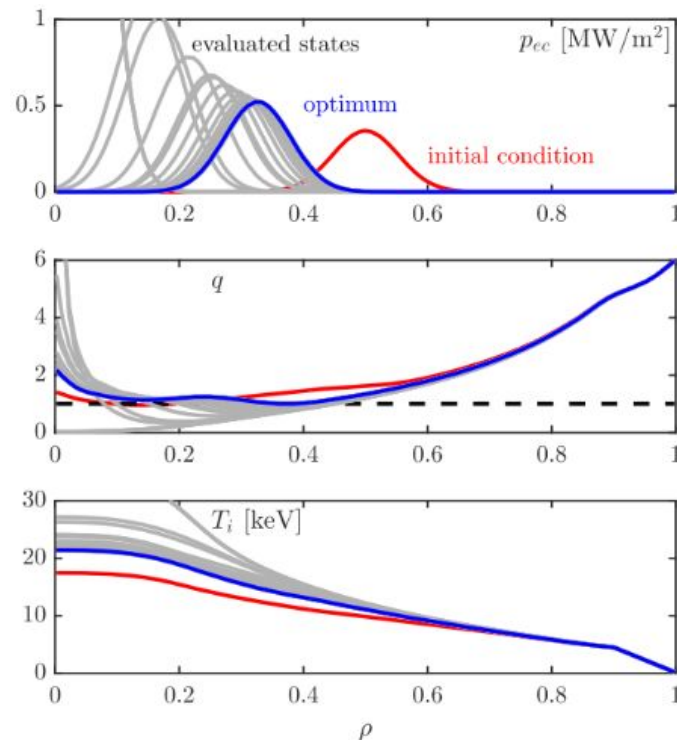
# Previous state-of-the-art in tokamak differentiable simulation

**RAPTOR** control-oriented simulator [Felici PPCF 2012]

MATLAB code, not-auto-differentiable, challenges with scalability and extensibility particularly with ML-surrogates

TORAX auto-differentiable with JAX:

- Easy to extend with any analytical models or ML-surrogate
- Easy to extend simulation sensitivity analysis and optimization with new parameter inputs
- ML-surrogate greatly facilitated by JAX's inherent support for neural network development and inference
- Python facilitates coupling within various workflows
- Scalable for large-scale validation and UQ



van Mulders NF 2021

SQP optimization of an ITER scenario

# JAX enables fast compiled and differentiable simulation with NumPy-like API

- JAX, originally developed by Google, is "NumPy on the CPU, GPU, TPU" (<https://jax.readthedocs.io/>)
  - Originally developed for AI/ML. Increasing applied for scientific computing
- Uses an updated version of Autograd to automatically differentiate NumPy-like code
  - Automatic transformation of functions to their *analytic* derivatives
- Uses TensorFlow's XLA (Accelerated Linear Algebra) JIT (just in time) compiler for speed
  - Same code can run seamlessly on CPU or accelerators



# Simple example of JAX function transformation



```
def sum_logistic(x):  
    return jnp.sum(1.0 / (1.0 + jnp.exp(-x)))  
  
x_small = jnp.arange(3.)  
derivative_fn = grad(sum_logistic)  
print(derivative_fn(x_small))
```

```
[0.25      0.19661194 0.10499357]
```

- Functions can be transformed into just-in-time (jit) compiled versions
  - e.g. `jitted_sum_logistic = jax.jit(sum_logistic)`
- Compile time overheads; for TORAX  $O(\text{compile-time}) \approx O(\text{run-time})$ 
  - Mitigated by compilation caching on memory or persistent (disk)
- In TORAX, bottleneck functions are JAX, e.g. PDE residual calculation + its Jacobian.
  - Glue code, control-flow, pre+post processing is standard Python.  
Eases development + coupling to wider frameworks.

# Governing equations: set of 1D flux-surface-averaged **nonlinear** transport PDEs

## Ion and electron heat equations

$$\frac{3}{2}V'^{-5/3} \left( \frac{\partial}{\partial t} - \frac{\dot{\Phi}_b}{2\Phi_b} \frac{\partial}{\partial \hat{\rho}} \right) [V'^{5/3} n_i T_i] =$$
$$\frac{1}{V'} \frac{\partial}{\partial \hat{\rho}} \left[ \chi_i n_i \frac{g_1}{V'} \frac{\partial T_i}{\partial \hat{\rho}} - g_0 q_i^{\text{conv}} T_i \right] + Q_i$$

$$\frac{3}{2}V'^{-5/3} \left( \frac{\partial}{\partial t} - \frac{\dot{\Phi}_b}{2\Phi_b} \frac{\partial}{\partial \hat{\rho}} \right) [V'^{5/3} n_e T_e] =$$
$$\frac{1}{V'} \frac{\partial}{\partial \hat{\rho}} \left[ \chi_e n_e \frac{g_1}{V'} \frac{\partial T_e}{\partial \hat{\rho}} - g_0 q_e^{\text{conv}} T_e \right] + Q_e$$

## Electron density equation


$$\left( \frac{\partial}{\partial t} - \frac{\dot{\Phi}_b}{2\Phi_b} \frac{\partial}{\partial \hat{\rho}} \right) [n_e V'] =$$
$$\frac{\partial}{\partial \hat{\rho}} \left[ D_e n_e \frac{g_1}{V'} \frac{\partial n_e}{\partial \hat{\rho}} - g_0 V_e n_e \right] + V' S_n$$

## Current diffusion equation

$$\frac{16\pi^2 \sigma_{||} \mu_0 \hat{\rho} \Phi_b^2}{F^2} \left( \frac{\partial \psi}{\partial t} - \frac{\dot{\rho} \dot{\Phi}_b}{2\Phi_b} \frac{\partial \psi}{\partial \hat{\rho}} \right) =$$
$$\frac{\partial}{\partial \hat{\rho}} \left( \frac{g_2 g_3}{\hat{\rho}} \frac{\partial \psi}{\partial \hat{\rho}} \right) - \frac{8\pi^2 V' \mu_0 \Phi_b}{F^2} \langle j_{ni} \rangle$$

- Moment equations of underlying kinetic equations
  - Toroidal symmetry + flux-surface averaging reduces to 1D
  - Scale-separation of turbulence, sources: "function calls" of reduced order models verified against high-fidelity
- Zero-gradient boundary conditions on axis
- Temperature, density
  - Dirichlet boundary conditions at edge
- $\psi$  (poloidal magnetic flux):
  - Neumann boundary condition at edge (fixed current)
  - or Dirichlet boundary condition for next timestep (edge voltage)
- Any subset of these equations are evolved: non-evolved profiles and initial conditions set in config with user data (xarray, numpy, python primitives)

# Spatial discretization: finite-volume-method with bespoke JAX fvm package



- TORAX JAX fvm package inspired by FiPy<sup>1</sup> and uses similar API
  - Constrained to solving convection-diffusion equations
- For (particle) convection, power-law  $\alpha$ -weighting scheme based on Péclet number
- Constructs nonlinear/linear systems of equations for solvers based on solver method

<sup>1</sup><https://www.ctcms.nist.gov/fipy/>



# Temporal discretization and system composition

Theta method, first-order in time.

$\theta=1$ , fully-implicit (default)

$$x_{t+\Delta t} - x_t = \Delta t [\theta F(x_{t+\Delta t}, t + \Delta t) + (1 - \theta) F(x_t, t)]$$

Governing equations more generally decomposed as follows

$$\frac{3}{2} V'^{-5/3} \left( \frac{\partial}{\partial t} - \frac{\dot{\Phi}_b}{2\Phi_b} \frac{\partial}{\partial \hat{\rho}} \right) [V'^{5/3} n_i T_i] = \frac{1}{V'} \frac{\partial}{\partial \hat{\rho}} \left[ \chi_i n_i \frac{g_1}{V'} \frac{\partial T_i}{\partial \hat{\rho}} - g_0 q_i^{\text{conv}} T_i \right] + Q_i \longrightarrow \tilde{\mathbf{T}}(x_{t+\Delta t}, u_{t+\Delta t}) \odot \mathbf{x}_{t+\Delta t} - \tilde{\mathbf{T}}(x_t, u_t) \odot \mathbf{x}_t = \Delta t [\theta (\bar{\mathbf{C}}(x_{t+\Delta t}, u_{t+\Delta t}) \mathbf{x}_{t+\Delta t} + \mathbf{c}(x_{t+\Delta t}, u_{t+\Delta t})) + (1 - \theta) (\bar{\mathbf{C}}(x_t, u_t) \mathbf{x}_t + \mathbf{c}(x_t, u_t))] ]$$

- $\mathbf{x}$  is state vector, subset of  $\{T_i, T_e, n_e, \psi\}$ , at time  $t$  or  $t+\Delta t$
- $\tilde{\mathbf{T}}$  is the "transient term", e.g.  $V'^{5/3} n_i$  in ion heat equation
- $\bar{\mathbf{C}}$  is the discretization matrix, including (possibly state-dependent) physics quantities like transport coefficients, geometry, etc.
- $\mathbf{c}$  is vector with source terms + boundary condition terms
- $\mathbf{u}$  corresponds to all "known" or quantities at time  $t$  and  $t+\Delta t$ , e.g. boundary conditions, heating amplitude trajectories, prescribed profiles, etc.

## Solver methods

$$\begin{aligned} \tilde{\mathbf{T}}(x_{t+\Delta t}, u_{t+\Delta t}) \odot \mathbf{x}_{t+\Delta t} - \tilde{\mathbf{T}}(x_t, u_t) \odot \mathbf{x}_t = \\ \Delta t \left[ \theta \left( \bar{\mathbf{C}}(x_{t+\Delta t}, u_{t+\Delta t}) \mathbf{x}_{t+\Delta t} + \mathbf{c}(x_{t+\Delta t}, u_{t+\Delta t}) \right) \right. \\ \left. + (1 - \theta) \left( \bar{\mathbf{C}}(x_t, u_t) \mathbf{x}_t + \mathbf{c}(x_t, u_t) \right) \right] \end{aligned} \quad (1)$$

- Predictor corrector: fixed point iteration for k (user-defined) steps on  $\mathbf{x}_{t+\Delta t}$  in  $\tilde{\mathbf{T}}, \bar{\mathbf{C}}, \mathbf{c}$ , starting from initial guess  $\mathbf{x}_t$

- Newton-Raphson nonlinear solver: iterative root finding for residual

$$\mathbf{R}(\mathbf{x}_{t+\Delta t}, \mathbf{x}_t, \mathbf{u}_{t+\Delta t}, \mathbf{u}_t, \theta, \Delta t) = 0$$

Where  $\mathbf{R}$  is the LHS-RHS of (1)

# Newton-Raphson illustration: simple example with heat diffusion

$$\frac{\partial T_k}{\partial t} = \frac{\partial}{\partial r} \left( \chi(T_{k+1}) \frac{\partial T_{k+1}}{\partial r} \right) + S(T_{k+1})$$

Fully implicit simple nonlinear diffusion equation

$$\frac{\vec{T}_{k+1} - \vec{T}_k}{\Delta t} - \bar{A}(\chi(T_{k+1}))\vec{T}_{k+1} - \vec{S}(T_{k+1}) \equiv \vec{R}(T_{k+1}, T_k) = 0$$

Discretize and define nonlinear system of equations to be solved (residual)

$$\vec{R}(T_{old}, T_k) + \bar{J}|_{(\vec{T}_{old})} (\vec{T}_{new} - \vec{T}_{old}) = 0$$

Newton-Raphson: starting from initial guess of  $T_{old}$  (e.g.  $T_k$ ), iteratively solve linear system for  $T_{new}$ , until  $R(T_{new}, T_k)$  within tolerance


All the physics goes into the residual function, and the Jacobian is JAX magic ✨

**jacobian = jax.jacfwd(residual)**

- TORAX has simple linesearch to ensure good Newton steps (physical  $T_{new}$ ), as well as  $\Delta t$  backtracking if no convergence

# Presently implemented physics models/couplings (non-exhaustive list)

- ML-surrogates where available
  - Turbulence (QLK-NN [van de Plassche 2020, Hamel 2025\*], TGLF-NN [Zanisi 2025])
  - Heating (TORIC-NN [Wallace APS 24])
- Fast analytical models where appropriate
  - Bootstrap current, neoclassical transport, ECCD
  - Mavrin polynomial fits to ADAS for line radiation, Bremsstrahlung
  - Fusion power, ion-electron collisional exchange, Ohmic power
- Pre-calculated sequence of geometry inputs: wrappers for CHEASE, FBT, EQDSK. IMAS underway
- Collaborations key
  - TORAX aims to be natural target platform for community-wide ML-surrogates
    - ONNX data storage for portability
  - Design focus on modularity



\* [https://github.com/google-deepmind/fusion\\_transport\\_surrogates](https://github.com/google-deepmind/fusion_transport_surrogates) (released last Friday)

# Runtimes and solver comparison

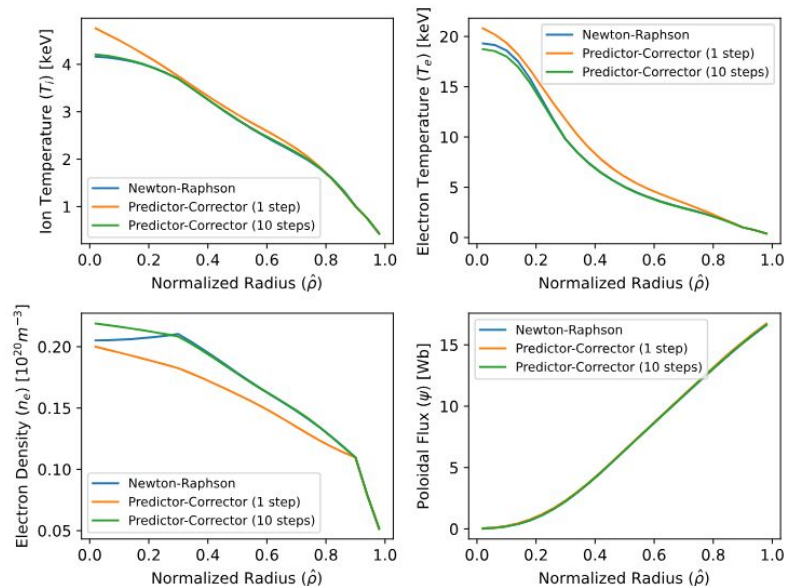


TABLE I. TORAX Simulation Performance Comparison

Solver	Compile [s]	Runtime [s]
Newton-Raphson	15.6	22
Predictor-Corrector (1 step)	4.5	6.5
Predictor-Corrector (10 steps)	4.6	8

Runtime and solver comparisons for a 80s ITER scenario current ramp-up.

- Faster than realtime for this config
- Still some low hanging fruit for runtime performance optimization

FIG. 5. Comparison of simulated  $T_i$ ,  $T_e$ ,  $n_e$ , and  $\psi$  profiles at  $t = 10$  s for the Newton-Raphson, Predictor-Corrector (1 step), and Predictor-Corrector (10 steps) solvers, using the example `iter-hybrid_rampup.py` configuration.

# Comparison with other integrated modelling suites

	JETTO/ASTRA	RAPTOR	TORAX
Coding language	Fortran	MATLAB	Python w/JAX
Linear solver + predictor-corrector	Yes	No	Yes
Newton-Raphson nonlinear solver	No	Yes	Yes
Discretization	FVM	FEM	FVM
Time-stepping	Adaptive	Deterministic	Adaptive or deterministic
Differentiable	No	Yes: manual	Yes: automatic
ML-surrogate coupling	Bespoke Fortran interface or <a href="https://github.com/Cambridge-ICCS/FTorch">https://github.com/Cambridge-ICCS/FTorch</a>	Mex-files from Fortran	Python w/JAX
Range of physics models	Broad, high-fidelity, no restriction (apart from speed)	Parameterized: formulas, ML-surrogates	<p>Any: analytical models + ML-surrogates can be in JAX kernel. Non-JAX models can be injected into kernel as explicit terms</p> <p>High-fidelity models can still be coupled into solver kernel: JAX compilation can be disabled. Allows evaluation of ML-surrogates against ground truth within the same framework</p>



# Towards full-sim differentiation with Forward Sensitivity Analysis: key is the Jacobian of the PDE system residual

State  $x_{k+1}$  solves residual at time  $k$ , e.g. with iterative Newton method.  
 $u$  is control input, parameterized by  $p$  (e.g. ECCD power waveform)

$$\tilde{f}_k \equiv \tilde{f}(x_{k+1}, x_k, u_k) = 0 \quad \forall k$$

Forward Sensitivity Analysis method


$$0 = \frac{d\tilde{f}_k}{dp} = \frac{\partial \tilde{f}_k}{\partial x_{k+1}} \frac{\partial x_{k+1}}{\partial p} + \frac{\partial \tilde{f}_k}{\partial x_k} \frac{\partial x_k}{\partial p} + \frac{\partial \tilde{f}_k}{\partial u_k} \frac{\partial u_k}{\partial p} + \frac{\partial \tilde{f}_k}{\partial p}$$

We want  $\frac{\partial x_{k+1}}{\partial p}$ , how the solution changes with respect to a control input modification.

Linear system above is recursively solved starting from initial condition  $\frac{\partial x_0}{\partial p}$

All  $f$  derivatives known and come from autodiff!

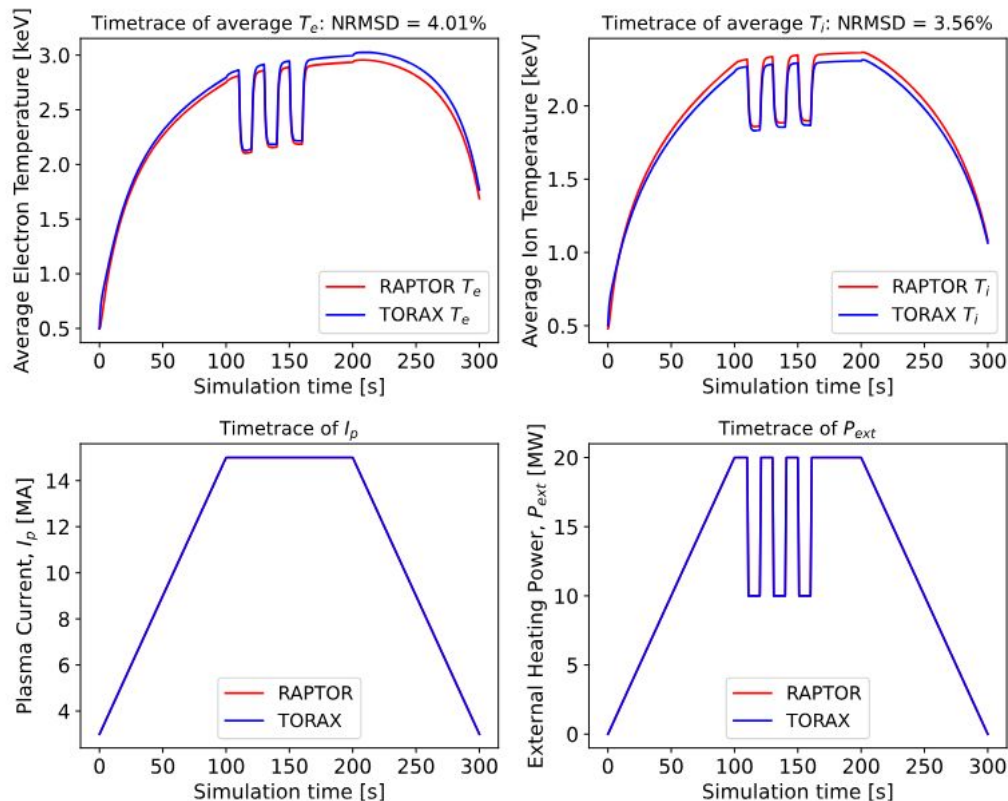
Key tool for sensitivity analysis, data-driven parameter identification, trajectory optimization methods



# Benchmarks vs RAPTOR

3

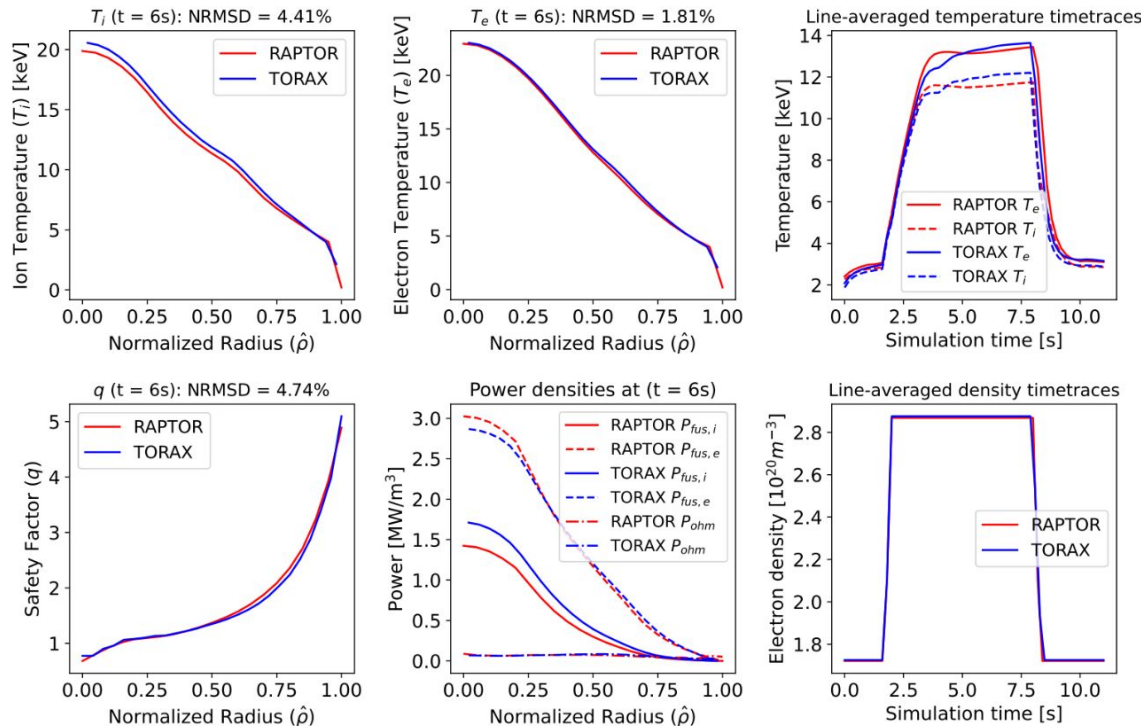
# Verification: TORAX vs RAPTOR agreement for ITER-like cases



## Modeling settings:

- ITER inspired params
- Nonlinear Newton-Raphson solver
- Heat transport + current diffusion
- 20MW heating: modulated
- Constant transport coefficients

# Verification: TORAX vs RAPTOR agreement on SPARC H-mode scenario



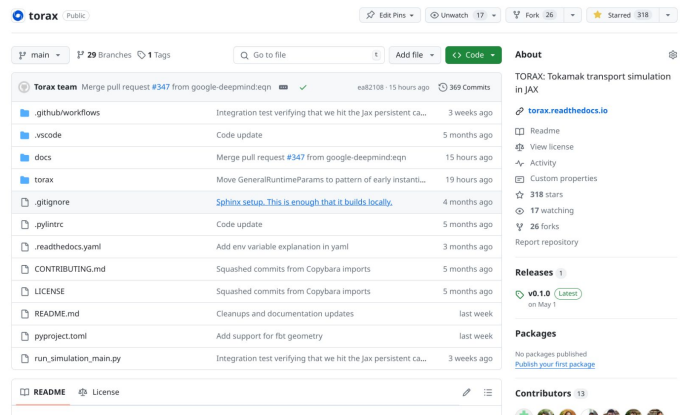
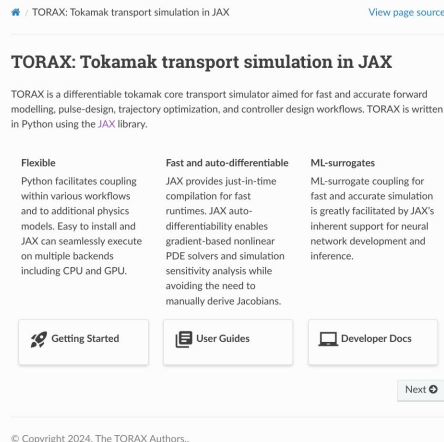
## Modeling settings:

- SPARC full-pulse scenario
- Nonlinear Newton-Raphson solver
- Heat transport + current diffusion
- Sequence of magnetic equilibria
- 11MW heating power
- ML-surrogate turbulent transport model (QLKNN10D\*)

$\Delta t = 0.2\text{s}$ : RAPTOR walltime: ~70s , TORAX walltime: ~14s

# TORAX open-source with permissive Apache 2.0 license

- Open source launch in June 2024
  - <https://github.com/google-deepmind/torax>
- Technical report + online documentation
  - [torax.readthedocs.io](https://torax.readthedocs.io)
  - <https://arxiv.org/abs/2406.06718>



arXiv > physics > arXiv:2406.06718

Physics > Plasma Physics

[Submitted on 10 Jun 2024 (v1), last revised 12 Jun 2024 (this version, v2)]

## TORAX: A Fast and Differentiable Tokamak Transport Simulator in JAX

Jonathan Citrin, Ian Goodfellow, Akhil Raju, Jeremy Chen, Jonas Degraeve, Craig Donner, Federico Felici, Philippe Hamel, Andrea Huber, Dmitry Nikulin, David Pfau, Brendan Tracey, Martin Riedmiller, Pushmeet Kohli

We present TORAX, a new, open-source, differentiable tokamak core transport simulator implemented in Python using the JAX framework. TORAX solves the coupled equations for ion heat transport, electron heat transport, particle transport, and current incorporating modular physics-based and ML models. JAX's just-in-time compilation ensures fast runtimes, while its automatic differentiation capability enables gradient-based optimization workflows and simplifies the use of Jacobian-based PDE solver: surrogates of physics models is greatly facilitated by JAX's intrinsic support for neural network development and inference. TORAX is verified against the established RAPTOR code, demonstrating agreement in simulated plasma profiles. TORAX provides a versatile tool for accelerating research in tokamak scenario modeling, pulse design, and control.

Comments: 16 pages, 7 figures

Subjects: Plasma Physics (physics.plasm-ph)

Cite as: arXiv:2406.06718 [physics.plasm-ph]

(or arXiv:2406.06718v2 [physics.plasm-ph] for this version)

<https://doi.org/10.48550/arXiv.2406.06718>

# Outlook

# 4

# Roadmap: towards higher physics fidelity and pulse planning applications

## Short term developments

- Reach physics feature parity and beyond compared to SOTA control-oriented tokamak sim
  - Sawteeth
  - Rotation + ExB shear
- Initial applications for tokamak pulse planning with collaborators

## Priorities for improved physics + numerics

- Multi-ion + impurity transport
- Improved ML-surrogates
  - Turbulence
  - Pedestal
  - Magnetic equilibrium
  - Heat/**particle sources**
  - **Plasma edge + wall**
- Demonstrate gradient-driven optimization applications
- Demonstrate large-scale batch simulations on GPU for optimization and UQ

## Validation/calibration against data

- Open source datasets for integrated modelling validation
  - Engagement with fusion community
- Data can also be used to further calibrate models, using gradient-driven methods



# Roadmap: initial applications and collaborations underway

## **CFS** MOSAIC Pulse Planning Workflow:

- GSPulse [Wai APS24]. Optimizes coil currents over full target pulse trajectory
- TORAX for internal plasma dynamics

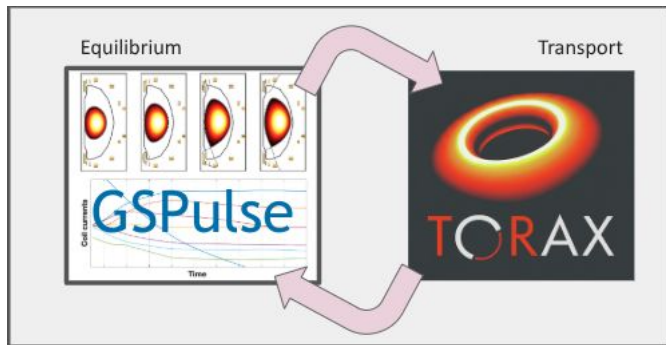


Image from Devon Battaglia

## **ITER** Pulse Design Simulator w/CEA + Ignition Computing

- Coupling to IMAS underway [Schneider, Bellourd, Sanders, van Vugt]

## **UKAEA**

- Progress in preparing TORAX for STEP simulation and benchmarking with JETTO [T. Brown, L. Zanisi]

## **CEA**

- E. Stancar for WEST trajectory optimization

# Summary

- TORAX, a new Python-based tokamak core transport code
  - With JAX: Fast for many-query applications, and auto-differentiable
  - Targeted for easy coupling to range of ML-surrogates
  - Verified against previous SOTA tokamak control-oriented simulator, RAPTOR
- Open-source for wider community impact
  - Supporting applications
  - Supporting integration of new physics models and ML-surrogates
  - Open source data for validation + calibration; keen to engage community on this effort
- Can couple to broader fusion simulation frameworks
  - Speed supports various applications in forward and inverse modelling
- Excited to see TORAX in action!



<https://arxiv.org/abs/2406.06718>  
<https://github.com/google-deepmind/torax>  
<https://torax.readthedocs.io>